vibrantlabsai / ragas Public

Code | Issues 303 | Pull requests 69 | Actions | Security and quality

3 Branches | 90 Tags | Go to file | Go to file | Code

veeceey fix: allow fork contributors in check-docs CI workflow (#2606)

298b682 · 2 months ago

.claude/commands	fix: metric inheritance patterns...	7 months ago
.cursor	docs: fix modifying-prompts-...	4 months ago
.github	fix: allow fork contributors in c...	2 months ago
docs	chore: remove office hours cal ...	3 months ago
examples	[AG-UI] Replacing "OpenAI" wi...	4 months ago
scripts	Retire `ragas.experimental` na...	8 months ago
src/ragas	fix: fix testset generation bugs ...	3 months ago
tests	perf: optimize average precisio...	3 months ago
.dockerignore	Fix/docs upgrade shade 2 (#2...	10 months ago
.gitattributes	chore: removed lfs files (#56)	3 years ago
.gitignore	feat: Add reusable testing infra...	6 months ago
.pre-commit-config.yaml	Add python 3.13 target, refact...	8 months ago
.readthedocs.yml	Merge workspace.toml into py...	8 months ago
CLAUDE.md	chore: cleanup old patterns an...	5 months ago
CODE_OF_CONDUCT.md	chore: add COC (#2437)	5 months ago
CONTRIBUTING.md	Feat: document PDF export wo...	4 months ago
LICENSE	chore: rebranding efforts (#24...	5 months ago
Makefile	feat(docs): add offline mermai...	4 months ago
README.md	chore: remove office hours cal ...	3 months ago

SECURITY.md	chore: rebranding efforts (#24...	5 months ago
mkdocs-pdf.yml	feat(docs): add offline mermai...	4 months ago
mkdocs.yml	feat(docs): add copy-to-llm bu...	4 months ago
pyproject.toml	feat(docs): add copy-to-llm bu...	4 months ago

ragas

Supercharge Your LLM Application Evaluations 🚀

release v0.4.3 | Made with Python | license Apache-2.0 | downloads/month 1M | chat 148 online | Ask DeepWiki

Documentation | **Quick start** | **Join Discord** | **Blog** | **NewsLetter** | **Careers**

Objective metrics, intelligent test generation, and data-driven insights for LLM apps

Ragas is your ultimate toolkit for evaluating and optimizing Large Language Model (LLM) applications. Say goodbye to time-consuming, subjective assessments and hello to data-driven, efficient evaluation workflows. Don't have a test dataset ready? We also do production-aligned test set generation.

Key Features

- 🎯 Objective Metrics: Evaluate your LLM applications with precision using both LLM-based and traditional metrics.
- 🧪 Test Data Generation: Automatically create comprehensive test datasets covering a wide range of scenarios.
- 🔗 Seamless Integrations: Works flawlessly with popular LLM frameworks like LangChain and major observability tools.
- 📊 Build feedback loops: Leverage production data to continually improve your LLM applications.

🛡️ Installation

Pypi:

```
pip install ragas
```

Alternatively, from source:

```
pip install git+https://github.com/vibrantlabsai/ragas
```

🔥 Quickstart

Clone a Complete Example Project

The fastest way to get started is to use the `ragas quickstart` command:

```
# List available templates
ragas quickstart

# Create a RAG evaluation project
ragas quickstart rag_eval

# Specify where you want to create it.
ragas quickstart rag_eval -o ./my-project
```

Available templates:

- `rag_eval` - Evaluate RAG systems

Coming Soon:

- `agent_evals` - Evaluate AI agents
- `benchmark_llm` - Benchmark and compare LLMs
- `prompt_evals` - Evaluate prompt variations
- `workflow_eval` - Evaluate complex workflows

Evaluate your LLM App

`ragas` comes with pre-built metrics for common evaluation tasks. For example, Aspect Critique evaluates any aspect of your output using `DiscreteMetric`:

```
import asyncio
from openai import AsyncOpenAI
from ragas.metrics import DiscreteMetric
```

```
from ragas.llms import llm_factory

# Setup your LLM
client = AsyncOpenAI()
llm = llm_factory("gpt-4o", client=client)

# Create a custom aspect evaluator
metric = DiscreteMetric(
    name="summary_accuracy",
    allowed_values=["accurate", "inaccurate"],
    prompt="""Evaluate if the summary is accurate and captures key information.

Response: {response}

Answer with only 'accurate' or 'inaccurate'."""
)

# Score your application's output
async def main():
    score = await metric.ascore(
        llm=llm,
        response="The summary of the text is..."
    )
    print(f"Score: {score.value}")  # 'accurate' or 'inaccurate'
    print(f"Reason: {score.reason}")


if __name__ == "__main__":
    asyncio.run(main())
```

> **Note**: Make sure your `OPENAI_API_KEY` environment variable is set.

Find the complete Quickstart Guide

Want help in improving your AI application using evals?

In the past 2 years, we have seen and helped improve many AI applications using evals. If you want help with improving and scaling up your AI application using evals.

🔗 Book a slot or drop us a line: founders@vibrantlabs.com.

🫂 Community

If you want to get more involved with Ragas, check out our discord server. It's a fun community where we geek out about LLM, Retrieval, Production issues, and more.

Contributors

```
+-------------------------------------------------------------------------+
|     +-----------------------------------------------------------------+     |
|     | Developers: Those who built with `ragas`.                       |     |
|     | (You have `import ragas` somewhere in your project)             |     |
|     |     +-----------------------------------------------------+     |     |
|     |     | Contributors: Those who make `ragas` better.        |     |     |
|     |     | (You make PR to this repo)                          |     |     |
|     |     +-----------------------------------------------------+     |     |
|     +-----------------------------------------------------------------+     |
+-------------------------------------------------------------------------+
```

We welcome contributions from the community! Whether it's bug fixes, feature additions, or documentation improvements, your input is valuable.

1. Fork the repository
2. Create your feature branch (git checkout -b feature/AmazingFeature)
3. Commit your changes (git commit -m 'Add some AmazingFeature')
4. Push to the branch (git push origin feature/AmazingFeature)
5. Open a Pull Request

🔍 Open Analytics

At Ragas, we believe in transparency. We collect minimal, anonymized usage data to improve our product and guide our development efforts.

✅ No personal or company-identifying information

✅ Open-source data collection code

✅ Publicly available aggregated data

To opt-out, set the `RAGAS_DO_NOT_TRACK` environment variable to `true`.

Cite Us

```
@misc{ragas2024,
  author       = {VibrantLabs},
  title        = {Ragas: Supercharge Your LLM Application Evaluations},
  year         = {2024},
```

Releases 89

v0.4.3 Latest
on Jan 13

+ 88 releases

Contributors 274

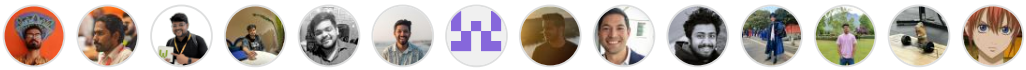

+ 260 contributors

Languages

Python 83.0% Jupyter Notebook 16.7% Other 0.3%