Claude API Docs

Models > Models overview

Models overview

Claude is a family of state-of-the-art large language models developed by Anthropic. This guide introduces the available models and compares their performance.

Copy page

Choosing a model

If you're unsure which model to use, consider starting with **Claude Opus 4.7** for the most complex tasks. It is our most capable generally available model, with a step-change improvement in agentic coding over Claude Opus 4.6.

All current Claude models support text and image input, text output, multilingual capabilities, and vision. Models are available via the Claude API, Amazon Bedrock, Vertex AI, and Microsoft Foundry.

Once you've picked a model, learn how to make your first API call.

Latest models comparison

Feature	Claude Opus 4.7
Description	Our most capable ge available model for c reasoning and agent
Claude API ID	`claude-opus-4-7`
Claude API alias	`claude-opus-4-7`

Configurações de cookies

Usamos cookies para fornecer e melhorar nossos serviços, analisar o uso do site e, se você concordar, personalizar sua experiência e comercializar nossos serviços para você. Você pode ler nossa Política de Cookies aqui.

Personalizar Configurações de Cookies

Rejeitar Todos os Cookies

Aceitar Todos os Cookies

Claude API Docs

AWS Bedrock ID	`7`[3]	`anthropic.claude-sonnet-4-6`	`haiku-4-5-20251001-v1:0`
GCP Vertex AI ID	`claude-opus-4-7`	`claude-sonnet-4-6`	`claude-haiku-4-5@20251001`
Pricing[1]	$5 / input MTok $25 / output MTok	$3 / input MTok $15 / output MTok	$1 / input MTok $5 / output MTok
Extended thinking	No	Yes	Yes
Adaptive thinking	Yes	Yes	No
Priority Tier	Yes	Yes	Yes
Comparative latency	Moderate	Fast	Fastest
Context window	1M tokens	1M tokens	200k tokens
Max output	128k tokens	64k tokens	64k tokens
Reliable knowledge cutoff	Jan 2026[2]	Aug 2025[2]	Feb 2025
Training data cutoff	Jan 2026	Jan 2026	Jul 2025

1 - See the pricing page for complete pricing inf
thinking costs, and vision processing fees.

*2 - **Reliable knowledge cutoff** indicates the dat*
***data cutoff** is the broader date range of trainin*

3 - Claude Opus 4.7 on AWS is available through

Claude Mythos Preview is offere
cybersecurity workflows as part
self-serve sign-up.

Configurações de cookies

Usamos cookies para fornecer e melhorar nossos serviços, analisar o uso do site e, se você concordar, personalizar sua experiência e comercializar nossos serviços para você. Você pode ler nossa Política de Cookies aqui.

Claude API Docs

rely on stable performance across different environments.

Starting with **Claude Sonnet 4.5 and all subsequent models** (including Claude Sonnet 4.6), AWS Bedrock offers two endpoint types: **global endpoints** (dynamic routing for maximum availability) and **regional endpoints** (guaranteed data routing through specific geographic regions). Google Vertex AI offers three endpoint types: global endpoints, **multi-region endpoints** (dynamic routing within a geographic area), and regional endpoints. For more information, see the third-party platform pricing section.

You can query model capabilities and token limits programmatically with the Models API. The response includes `max_input_tokens`, `max_tokens`, and a `capabilities` object for every available model.

The Max output values above apply to the synchronous Messages API. On the Message Batches API, Opus 4.7, Opus 4.6, and Sonnet 4.6 support up to 300k output tokens by using the `output-300k-2026-03-24` beta header.

Legacy models

Prompt and output performance

Claude 4 models excel in:

- **Performance**: Top-tier results i
  handling, honesty, and image p
  information.
- **Engaging responses**: Claude m
  like interactions.
  - If you prefer more concise
    model toward the desired
    details.
  - For prompting best practic

Configurações de cookies

Usamos cookies para fornecer e melhorar nossos serviços, analisar o uso do site e, se você concordar, personalizar sua experiência e comercializar nossos serviços para você. Você pode ler nossa Política de Cookies aqui.

Claude API Docs

Migrating to Claude Opus 4.7

If you're currently using Claude Opus 4.6 or older Claude models, consider migrating to Claude Opus 4.7 to take advantage of improved intelligence and a step-change jump in agentic coding. For detailed migration instructions, see Migrating to Claude Opus 4.7.

Get started with Claude

If you're ready to start exploring what Claude can do for you, dive in! Whether you're a developer looking to integrate Claude into your applications or a user wanting to experience the power of AI firsthand, the following resources can help.

Looking to chat with Claude? Visit claude.ai!

Intro to Claude

Explore Claude's capabilities and development flow.

Quickstart

Learn how to make your first API call i

Claude Console

Craft and test powerful prompts dire

If you have any questions or need a
or consult the Discord community.

Configurações de cookies

Usamos cookies para fornecer e melhorar nossos serviços, analisar o uso do site e, se você concordar, personalizar sua experiência e comercializar nossos serviços para você. Você pode ler nossa Política de Cookies aqui.

Claude API Docs

Claude API Docs

Solutions

AI agents

Code modernization

Coding

Customer support

Education

Financial services

Government

Life sciences

Partners

Amazon Bedrock

Google Cloud's Vertex AI

Learn

Blog

Courses

Use cases

Connectors

Customer stories

Engineering at Anthropic

Events

Powered by Claude

Service partners

Startups program

Company

Anthropic

Help and security

Availability

Status

Configurações de cookies

Usamos cookies para fornecer e melhorar nossos serviços, analisar o uso do site e, se você concordar, personalizar sua experiência e comercializar nossos serviços para você. Você pode ler nossa Política de Cookies aqui.

Claude API Docs

Terms and policies

Privacy policy

Responsible disclosure policy

Terms of service: Commercial

Terms of service: Consumer

Usage policy

Configurações de cookies

Usamos cookies para fornecer e melhorar nossos serviços, analisar o uso do site e, se você concordar, personalizar sua experiência e comercializar nossos serviços para você. Você pode ler nossa Política de Cookies aqui.